

Deep Learning Software

PyTorch

Department of Computer Science, NCTU

TA Yu-Chuan Chuang

Some slides are from Stanford CS231n

Frameworks



Caffe

And others.....

Frameworks



Most people use these



Caffe

And others.....

Frameworks

We will focus on this



And others.....

The advantages of deep learning frameworks

- Developing and testing new ideas are quickly.
- Computing gradients automatically
- Running model structures on GPU is efficiently.

Please use **PyTorch to complete all your assignments !!**

 PyTorch  PyTorch  PyTorch

Install PyTorch

- <http://pytorch.org/>
- <https://pytorch.org/get-started/previous-versions/>
- <https://www.anaconda.com/>




Get Started

TA's environment

PyTorch Fundamental Concepts

- **Tensor** : Like a numpy array, but can run on GPU.



- **Autograd** : Package for building computational graphs out of Tensors and automatically computing gradients.
- **Module** : A neural network layer; may store state or learnable weights.

Computational Graphs


Numpy

```
import numpy as np
np.random.seed(0)

N, D = 3, 4

x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)

a = x * y
b = a + z
c = np.sum(b)
```



Neural network can be denoted as a directed acyclic graph

Computational Graphs

Numpy

```
import numpy as np
np.random.seed(0)
```


```
N, D = 3, 4
```

```
x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)
```

```
a = x * y
b = a + z
c = np.sum(b)
```

```
grad_c = 1.0
grad_b = grad_c * np.ones((N, D))
grad_a = grad_b.copy()
grad_z = grad_b.copy()
grad_x = grad_a * y
grad_y = grad_a * x
```

compute gradients



Problems:

- Can't run on GPU
- Have to compute our own gradients

Computational Graphs

Numpy


```
import numpy as np
np.random.seed(0)
```

```
N, D = 3, 4
```

```
x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)
```

```
a = x * y
b = a + z
c = np.sum(b)
```

```
grad_c = 1.0
grad_b = grad_c * np.ones((N, D))
grad_a = grad_b.copy()
grad_z = grad_b.copy()
grad_x = grad_a * y
grad_y = grad_a * x
```



PyTorch

```
import torch
```

```
N, D = 3, 4
```

```
x = torch.randn(N, D)
y = torch.randn(N, D)
z = torch.randn(N, D)
```

```
a = x * y
b = a + z
c = torch.sum(b)
```

Looks exactly like numpy!

Computational Graphs

Numpy


```
import numpy as np
np.random.seed(0)

N, D = 3, 4

x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)

a = x * y
b = a + z
c = np.sum(b)

grad_c = 1.0
grad_b = grad_c * np.ones((N, D))
grad_a = grad_b.copy()
grad_z = grad_b.copy()
grad_x = grad_a * y
grad_y = grad_a * x
```



PyTorch

```
import torch

N, D = 3, 4
x = torch.randn(N, D, requires_grad=True)
y = torch.randn(N, D)
z = torch.randn(N, D)

a = x * y
b = a + z
c = torch.sum(b)

c.backward()
print(x.grad)
```

PyTorch handles gradients for us!

.backward() compute gradient

Computational Graphs

Numpy


```
import numpy as np
np.random.seed(0)

N, D = 3, 4

x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)

a = x * y
b = a + z
c = np.sum(b)

grad_c = 1.0
grad_b = grad_c * np.ones((N, D))
grad_a = grad_b.copy()
grad_z = grad_b.copy()
grad_x = grad_a * y
grad_y = grad_a * x
```



PyTorch

```
import torch
device = 'cuda:0'
N, D = 3, 4
x = torch.randn(N, D, requires_grad=True,
                 device=device)
y = torch.randn(N, D, device=device)
z = torch.randn(N, D, device=device)

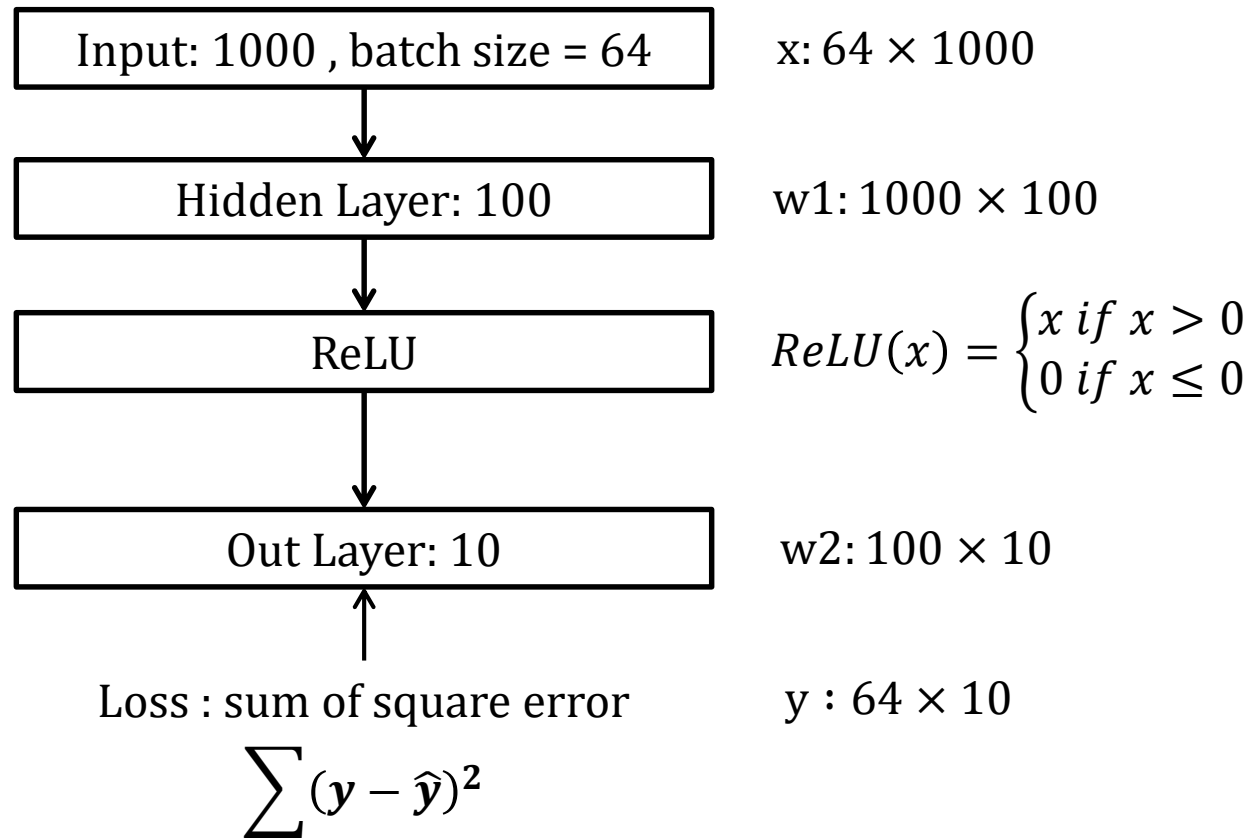
a = x * y
b = a + z
c = torch.sum(b)

c.backward()
print(x.grad)
```

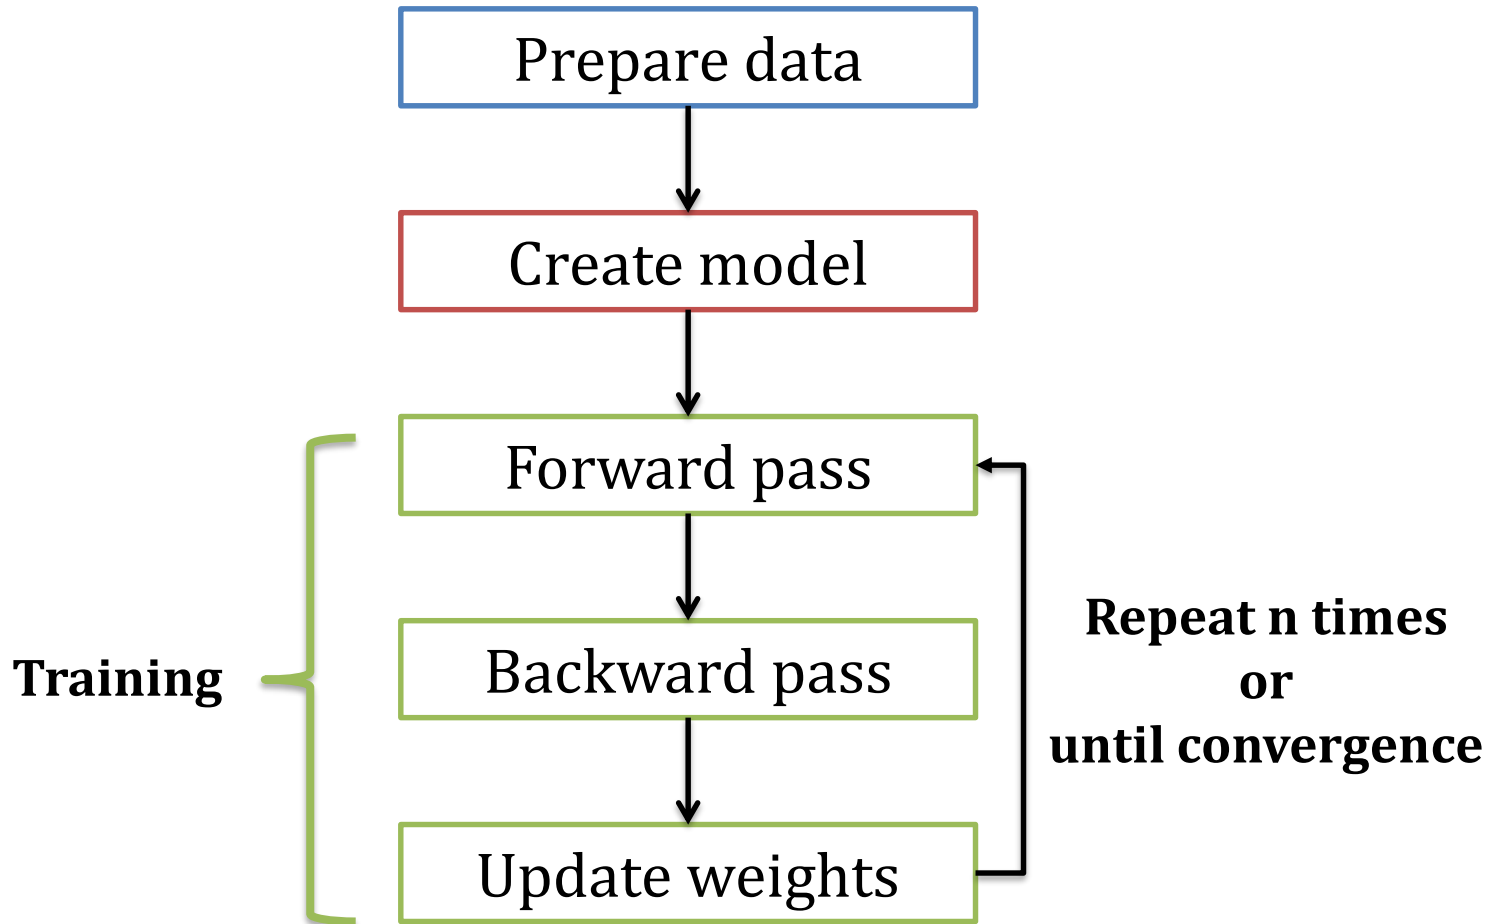
Trivial to run on GPU - just construct arrays on a different device!

Example

- 2-layer network



Flow Chart



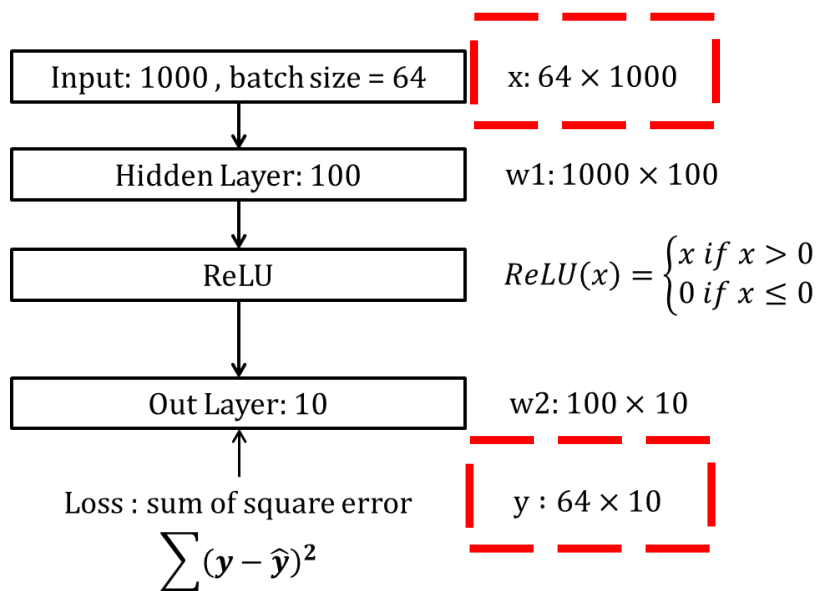
Step1. Prepare Data

PyTorch Tensors

Create random tensors as input and ground truth

To run on GPU, just use a different device, like a following:

```
device = torch.device('cuda:0')
```



```
import torch
device = torch.device('cpu')
learning_rate = 1e-6
```

```
x = torch.randn(64, 1000, device=device)
y = torch.randn(64, 10, device=device)
```

```
w1 = torch.randn(1000, 100, device=device)
w2 = torch.randn(100, 10, device=device)
```

```
for i in range(500):
    h = x.mm(w1)
    h_relu = h.clamp(min=0)
    y_pred = h_relu.mm(w2)
    loss = (y_pred - y).pow(2).sum()
```

```
grad_y_pred = 2.0 * (y_pred - y)
grad_h_relu = grad_y_pred.mm(w2.t())
grad_h = grad_h_relu.clone()
grad_h[h < 0] = 0
```

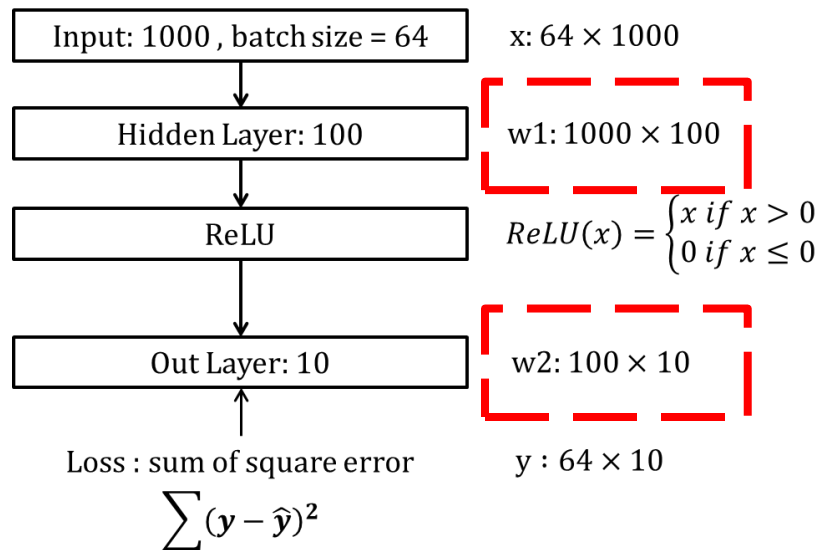
```
grad_w2 = h_relu.t().mm(grad_y_pred)
grad_w1 = x.t().mm(grad_h)
```

```
w1 -= learning_rate * grad_w1
w2 -= learning_rate * grad_w2
```


Step2. Create Model

PyTorch Tensors

Create random tensors as layer weights



```
import torch
device = torch.device('cpu')
learning_rate = 1e-6

x = torch.randn(64, 1000, device=device)
y = torch.randn(64, 10, device=device)

w1 = torch.randn(1000, 100, device=device)
w2 = torch.randn(100, 10, device=device)

for i in range(500):
    h = x.mm(w1)
    h_relu = h.clamp(min=0)
    y_pred = h_relu.mm(w2)
    loss = (y_pred - y).pow(2).sum()

    grad_y_pred = 2.0 * (y_pred - y)
    grad_h_relu = grad_y_pred.mm(w2.t())
    grad_h = grad_h_relu.clone()
    grad_h[h < 0] = 0

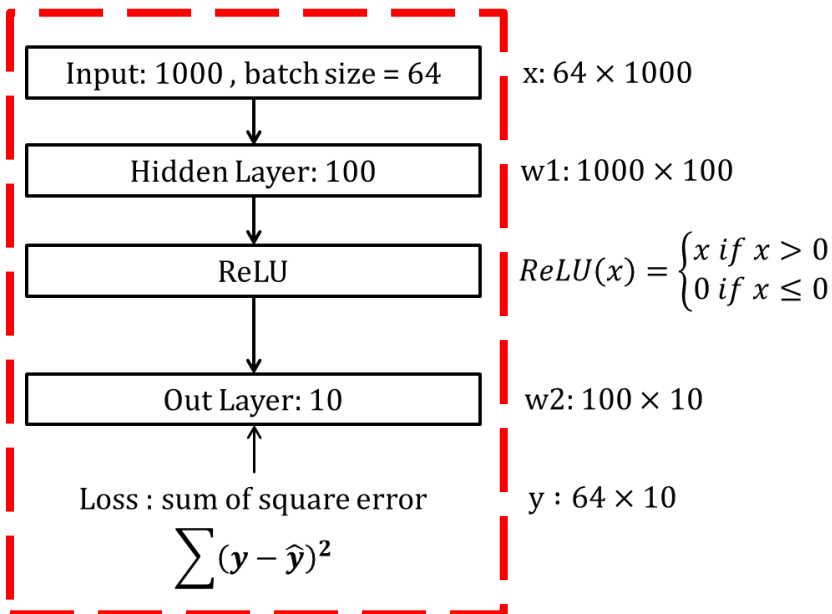
    grad_w2 = h_relu.t().mm(grad_y_pred)
    grad_w1 = x.t().mm(grad_h)

    w1 -= learning_rate * grad_w1
    w2 -= learning_rate * grad_w2
```


Step3. Forward pass

PyTorch Tensors

Compute predictions and loss



```
import torch
device = torch.device('cpu')
learning_rate = 1e-6

x = torch.randn(64, 1000, device=device)
y = torch.randn(64, 10, device=device)

w1 = torch.randn(1000, 100, device=device)
w2 = torch.randn(100, 10, device=device)

for i in range(500):
    h = x.mm(w1)
    h_relu = h.clamp(min=0)
    y_pred = h_relu.mm(w2)
    loss = (y_pred - y).pow(2).sum()

    grad_y_pred = 2.0 * (y_pred - y)
    grad_h_relu = grad_y_pred.mm(w2.t())
    grad_h = grad_h_relu.clone()
    grad_h[h < 0] = 0

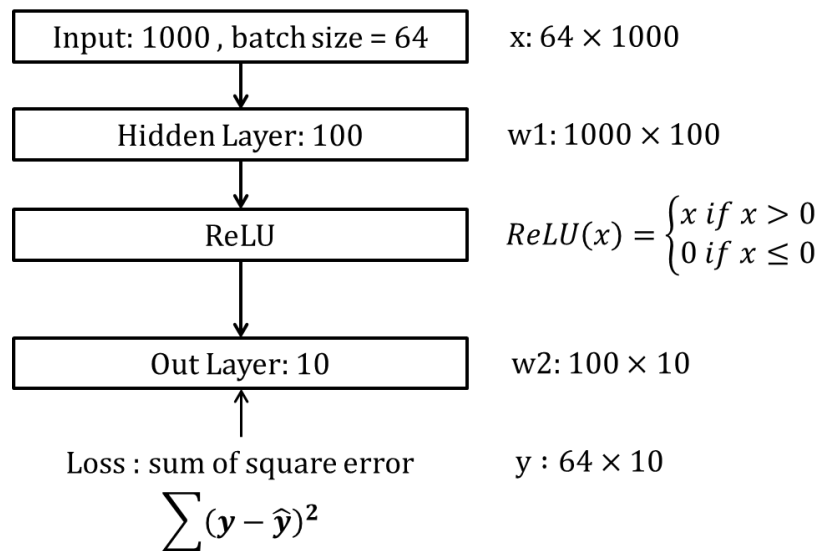
    grad_w2 = h_relu.t().mm(grad_y_pred)
    grad_w1 = x.t().mm(grad_h)

    w1 -= learning_rate * grad_w1
    w2 -= learning_rate * grad_w2
```

Step4. Backward pass

PyTorch Tensors

Manually compute gradients



```
import torch
device = torch.device('cpu')
learning_rate = 1e-6

x = torch.randn(64, 1000, device=device)
y = torch.randn(64, 10, device=device)

w1 = torch.randn(1000, 100, device=device)
w2 = torch.randn(100, 10, device=device)

for i in range(500):
    h = x.mm(w1)
    h_relu = h.clamp(min=0)
    y_pred = h_relu.mm(w2)
    loss = (y_pred - y).pow(2).sum()

    grad_y_pred = 2.0 * (y_pred - y)
    grad_h_relu = grad_y_pred.mm(w2.t())
    grad_h = grad_h_relu.clone()
    grad_h[h < 0] = 0


    grad_w2 = h_relu.t().mm(grad_y_pred)
    grad_w1 = x.t().mm(grad_h)

    w1 -= learning_rate * grad_w1
    w2 -= learning_rate * grad_w2
```

Step5. Update Weights

PyTorch Tensors

Gradient descent step on weights



```
import torch
device = torch.device('cpu')
learning_rate = 1e-6

x = torch.randn(64, 1000, device=device)
y = torch.randn(64, 10, device=device)

w1 = torch.randn(1000, 100, device=device)
w2 = torch.randn(100, 10, device=device)

for i in range(500):
    h = x.mm(w1)
    h_relu = h.clamp(min=0)
    y_pred = h_relu.mm(w2)
    loss = (y_pred - y).pow(2).sum()

    grad_y_pred = 2.0 * (y_pred - y)
    grad_h_relu = grad_y_pred.mm(w2.t())
    grad_h = grad_h_relu.clone()
    grad_h[h < 0] = 0

    grad_w2 = h_relu.t().mm(grad_y_pred)
    grad_w1 = x.t().mm(grad_h)

    w1 -= learning_rate * grad_w1
    w2 -= learning_rate * grad_w2
```

Easily implement your own
deep learning model by using
PyTorch

Step1. Prepare Data

PyTorch.utils.data

DataLoader wraps a **Dataset** and provides minibatching, shuffling, multithreading, for you.

When you need to load custom data, just write your own Dataset class.

Iterate over loader to form minibatches

<https://github.com/utkuozbulak/pytorch-custom-dataset-examples>

```
import torch
from torch.utils.data import TensorDataset, DataLoader

device = torch.device('cpu')
x = torch.randn(64, 1000)
y = torch.randn(64, 10)
loader = DataLoader(TensorDataset(x, y), batch_size=8)

class TwoLayerNet(torch.nn.Module):
    def __init__(self, D_in, H, D_out):
        super(TwoLayerNet, self).__init__()
        self.linear_1 = torch.nn.Linear(D_in, H)
        self.linear_2 = torch.nn.Linear(H, D_out)
        self.relu = torch.nn.ReLU()

    def forward(self, x):
        h = self.linear_1(x)
        h_relu = self.relu(h)
        y_pred = self.linear_2(h_relu)
        return y_pred

model = TwoLayerNet(D_in=1000, H=100, D_out=10)
optimizer = torch.optim.SGD(model.parameters(),
                              lr=1e-1)

for epochs in range(500):
    for x_batch, y_batch in loader:
        optimizer.zero_grad()
        y_pred = model(x_batch)
        loss = torch.nn.functional.mse_loss(y_pred,
                                              y_batch)

        loss.backward()
        optimizer.step()
        print(loss.item())
```

Step2. Create Model

PyTorch.nn

Higher-level wrapper for working with neural nets

Use this ! It will make your life easier.

A PyTorch Module is a neural net layer, it can contain weights or other modules.

Define your whole model as a single module.

```
import torch
from torch.utils.data import TensorDataset, DataLoader

device = torch.device('cpu')
x = torch.randn(64, 1000)
y = torch.randn(64, 10)
loader = DataLoader(TensorDataset(x, y), batch_size=8)
```

```
class TwoLayerNet(torch.nn.Module):
    def __init__(self, D_in, H, D_out):
        super(TwoLayerNet, self).__init__()
        self.linear_1 = torch.nn.Linear(D_in, H)
        self.linear_2 = torch.nn.Linear(H, D_out)
        self.relu = torch.nn.ReLU()

    def forward(self, x):
        h = self.linear_1(x)
        h_relu = self.relu(h)
        y_pred = self.linear_2(h_relu)
        return y_pred
```

```
model = TwoLayerNet(D_in=1000, H=100, D_out=10)
optimizer = torch.optim.SGD(model.parameters(),
                             lr=1e-1)

for epochs in range(500):
    for x_batch, y_batch in loader:
        optimizer.zero_grad()
        y_pred = model(x_batch)
        loss = torch.nn.functional.mse_loss(y_pred,
                                             y_batch)

        loss.backward()
        optimizer.step()
        print(loss.item())
```

Step2. Create Model

PyTorch.nn

Initializer sets up two children
(Module can contain Modules)

```
import torch
from torch.utils.data import TensorDataset, DataLoader

device = torch.device('cpu')
x = torch.randn(64, 1000)
y = torch.randn(64, 10)
loader = DataLoader(TensorDataset(x, y), batch_size=8)

class TwoLayerNet(torch.nn.Module):
    def __init__(self, D_in, H, D_out):
        super(TwoLayerNet, self).__init__()
        self.linear_1 = torch.nn.Linear(D_in, H)
        self.linear_2 = torch.nn.Linear(H, D_out)
        self.relu = torch.nn.ReLU()

    def forward(self, x):
        h = self.linear_1(x)
        h_relu = self.relu(h)
        y_pred = self.linear_2(h_relu)
        return y_pred

model = TwoLayerNet(D_in=1000, H=100, D_out=10)
optimizer = torch.optim.SGD(model.parameters(),
                              lr=1e-1)

for epochs in range(500):
    for x_batch, y_batch in loader:
        optimizer.zero_grad()
        y_pred = model(x_batch)
        loss = torch.nn.functional.mse_loss(y_pred,
                                              y_batch)

        loss.backward()
        optimizer.step()
        print(loss.item())
```

Step2. Create Model

PyTorch.nn

Define forward pass using child modules

No need to define backward – autograd will handle it.

```
import torch
from torch.utils.data import TensorDataset, DataLoader

device = torch.device('cpu')
x = torch.randn(64, 1000)
y = torch.randn(64, 10)
loader = DataLoader(TensorDataset(x, y), batch_size=8)

class TwoLayerNet(torch.nn.Module):
    def __init__(self, D_in, H, D_out):
        super(TwoLayerNet, self).__init__()
        self.linear_1 = torch.nn.Linear(D_in, H)
        self.linear_2 = torch.nn.Linear(H, D_out)
        self.relu = torch.nn.ReLU()

    def forward(self, x):
        h = self.linear_1(x)
        h_relu = self.relu(h)
        y_pred = self.linear_2(h_relu)
        return y_pred

model = TwoLayerNet(D_in=1000, H=100, D_out=10)
optimizer = torch.optim.SGD(model.parameters(),
                              lr=1e-1)

for epochs in range(500):
    for x_batch, y_batch in loader:
        optimizer.zero_grad()
        y_pred = model(x_batch)
        loss = torch.nn.functional.mse_loss(y_pred,
                                              y_batch)

        loss.backward()
        optimizer.step()
        print(loss.item())
```


Step3. Forward pass

PyTorch.nn

Define forward pass using child modules

Feed data to model, and compute loss

nn.functional has a lot of useful helpers like loss functions

```
import torch
from torch.utils.data import TensorDataset, DataLoader

device = torch.device('cpu')
x = torch.randn(64, 1000)
y = torch.randn(64, 10)
loader = DataLoader(TensorDataset(x, y), batch_size=8)

class TwoLayerNet(torch.nn.Module):
    def __init__(self, D_in, H, D_out):
        super(TwoLayerNet, self).__init__()
        self.linear_1 = torch.nn.Linear(D_in, H)
        self.linear_2 = torch.nn.Linear(H, D_out)
        self.relu = torch.nn.ReLU()

    def forward(self, x):
        h = self.linear_1(x)
        h_relu = self.relu(h)
        y_pred = self.linear_2(h_relu)
        return y_pred

model = TwoLayerNet(D_in=1000, H=100, D_out=10)
optimizer = torch.optim.SGD(model.parameters(),
                              lr=1e-1)

for epochs in range(500):
    for x_batch, y_batch in loader:
        optimizer.zero_grad()
        y_pred = model(x_batch)
        loss = torch.nn.functional.mse_loss(y_pred,
                                              y_batch)

        loss.backward()
        optimizer.step()
        print(loss.item())
```

Step4. Backward pass

PyTorch.autograd

Forward pass looks exactly the same as before, but we don't need to track intermediate values.

PyTorch keeps track of them for us in the computational graph.

Compute gradient of loss with respect to all model weights (they have `requires_grad=True`)

```
import torch
from torch.utils.data import TensorDataset, DataLoader

device = torch.device('cpu')
x = torch.randn(64, 1000)
y = torch.randn(64, 10)
loader = DataLoader(TensorDataset(x, y), batch_size=8)

class TwoLayerNet(torch.nn.Module):
    def __init__(self, D_in, H, D_out):
        super(TwoLayerNet, self).__init__()
        self.linear_1 = torch.nn.Linear(D_in, H)
        self.linear_2 = torch.nn.Linear(H, D_out)
        self.relu = torch.nn.ReLU()

    def forward(self, x):
        h = self.linear_1(x)
        h_relu = self.relu(h)
        y_pred = self.linear_2(h_relu)
        return y_pred

model = TwoLayerNet(D_in=1000, H=100, D_out=10)
optimizer = torch.optim.SGD(model.parameters(),
                              lr=1e-1)

for epochs in range(500):
    for x_batch, y_batch in loader:
        optimizer.zero_grad()
        y_pred = model(x_batch)
        loss = torch.nn.functional.mse_loss(y_pred,
                                              y_batch)
        loss.backward()
        optimizer.step()
        print(loss.item())
```

Step5. Update Weights

PyTorch.optim

Use an **optimizer** for different update rules.

After computing gradients, use optimizer to update each model parameters and reset gradients

```
import torch
from torch.utils.data import TensorDataset, DataLoader

device = torch.device('cpu')
x = torch.randn(64, 1000)
y = torch.randn(64, 10)
loader = DataLoader(TensorDataset(x, y), batch_size=8)

class TwoLayerNet(torch.nn.Module):
    def __init__(self, D_in, H, D_out):
        super(TwoLayerNet, self).__init__()
        self.linear_1 = torch.nn.Linear(D_in, H)
        self.linear_2 = torch.nn.Linear(H, D_out)
        self.relu = torch.nn.ReLU()

    def forward(self, x):
        h = self.linear_1(x)
        h_relu = self.relu(h)
        y_pred = self.linear_2(h_relu)
        return y_pred


model = TwoLayerNet(D_in=1000, H=100, D_out=10)
optimizer = torch.optim.SGD(model.parameters(), lr=1e-1)

for epochs in range(500):
    for x_batch, y_batch in loader:
        optimizer.zero_grad()
        y_pred = model(x_batch)
        loss = torch.nn.functional.mse_loss(y_pred, y_batch)

        loss.backward()
        optimizer.step()
        print(loss.item())
```

Real Application

- MNIST example for PyTorch



- `git clone https://github.com/JiaRenChang/DLcourse_NCTU.git`

 [CNN_MNIST_pytorch.py](#)

Build and train a CNN classifier

- Data Loader
- Define Network
- Define Optimizer/Loss function
- Learning rate scheduling
- Training
- Testing
- Run and Save model

Set hypermeters

```
# Training settings
parser = argparse.ArgumentParser(description='PyTorch MNIST Example')
parser.add_argument('--batch-size', type=int, default=128, metavar='N',
                    help='input batch size for training (default: 128)')
parser.add_argument('--epochs', type=int, default=20, metavar='N',
                    help='number of epochs to train (default: 164)')
parser.add_argument('--lr', type=float, default=0.1, metavar='LR',
                    help='learning rate (default: 0.1)')
parser.add_argument('--momentum', type=float, default=0.9, metavar='M',
                    help='SGD momentum (default: 0.9)')
parser.add_argument('--no-cuda', action='store_true', default=False,
                    help='disables CUDA training')
parser.add_argument('--seed', type=int, default=1, metavar='S',
                    help='random seed (default: 1)')

args = parser.parse_args()
```

Data Loader

- Pytorch offers data loaders for popular dataset

The following datasets are available:

Datasets

- MNIST
- COCO
 - Captions
 - Detection
- LSUN
- ImageFolder
- Imagenet-12
- CIFAR
- STL10
- SVHN
- PhotoTour

Data Loader


```
import torch
from torchvision import datasets, transforms

# Dataloader
train_loader = torch.utils.data.DataLoader(
    datasets.MNIST('../data', train=True, download=True,
                   transform=transforms.Compose([
                       transforms.ToTensor(),
                       transforms.Normalize((0.1307,), (0.3081,))
                   ])),
    batch_size=args.batch_size, shuffle=True, num_workers = 2)

test_loader = torch.utils.data.DataLoader(
    datasets.MNIST('../data', train=False, transform=transforms.Compose([
        transforms.ToTensor(),
        transforms.Normalize((0.1307,), (0.3081,))
    ])),
    batch_size=args.batch_size, shuffle=True, num_workers = 2)
```


Define Network

- LeNet



Define Network

```
#Define Network, we implement LeNet here
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 6, kernel_size=(5,5), stride=1, padding=0)
        self.conv2 = nn.Conv2d(6, 16, kernel_size=(5,5), stride=1, padding=0)
        self.fc1 = nn.Linear(16*4*4, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        out = F.relu(self.conv1(x))
        out = F.max_pool2d(out, 2)
        out = F.relu(self.conv2(out))
        out = F.max_pool2d(out, 2)
        out = out.view(out.size(0), -1) #flatten
        out = F.relu(self.fc1(out))
        out = F.relu(self.fc2(out))
        out = self.fc3(out)
        return out

model = Net()
if args.cuda:
    device = torch.device('cuda')
    model.to(device)
```

Define Optimizer/Loss function

- Cross Entropy Loss
- Stochastic Gradient Descent

```
#define optimizer/loss function  
Loss = nn.CrossEntropyLoss()  
optimizer = optim.SGD(model.parameters(), lr=args.lr, momentum=args.momentum)
```

Learning rate scheduling

- 20 epochs
- LR decay at 10 and 15 epoch

```
#learning rate scheduling
def adjust_learning_rate(optimizer, epoch):

    if epoch < 10:
        lr = 0.01
    elif epoch < 15:
        lr = 0.001
    else:
        lr = 0.0001

    for param_group in optimizer.param_groups:
        param_group['lr'] = lr
```

Training

```
#training function
def train(epoch):
    model.train() Set model to training mode
    adjust_learning_rate(optimizer, epoch)

    for batch_idx, (data, target) in enumerate(train_loader):
        if args.cuda:
            data, target = data.to(device), target.to(device)

            optimizer.zero_grad() Clean gradient
            output = model(data)
            loss = Loss(output, target)
            loss.backward() Backward gradient
            optimizer.step() Update weight
        print('Train Epoch: {} [{}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
            epoch, batch_idx * len(data), len(train_loader.dataset),
            100. * batch_idx / len(train_loader), loss.data[0]))
```

Testing

```
#Testing function
def test(epoch):
    model.eval()
    test_loss = 0
    correct = 0
    for batch_idx, (data, target) in enumerate(test_loader):
        if args.cuda:
            data, target = data.to(device), target.to(device)
        with torch.no_grad():
            output = model(data)
        test_loss += Loss(output, target).data[0]
        pred = output.data.max(1)[1] # get the index of the max log-probability
        correct += pred.eq(target.data).cpu().sum()

    test_loss = test_loss
    test_loss /= len(test_loader) # loss function already averages over batch size
    print('\nTest set: Average loss: {:.4f}, Accuracy: {}/{} ({:.0f}%)'.format(
        test_loss, correct, len(test_loader.dataset),
        100. * correct / len(test_loader.dataset)))
```

Run and Save model

```
#run and save model
for epoch in range(1, args.epochs + 1):
    train(epoch)
    test(epoch)
    savefilename = 'LeNet_'+str(epoch)+'.tar'
    torch.save({
        'epoch': epoch,
        'state_dict': model.state_dict(),
    }, savefilename)
```

You can achieve ~99.1% test accuracy.

Exercise

- Deeper: add more convolution layer
 - insert two 3x3 conv layer between conv1 and conv2 (stride=1,pad=1)
 - Hint: define new conv layer, and forward
 - Notice the *spatial* dimension
- Wider: add more neuron
 - Make your net 2x wider
 - Notice the *in/out* dimension
- Other Optimizer
 - Try Adam/RMSprop
- More epochs, New learning rate schedule,